

# HARSH DADHICH

San Francisco, CA | [me@harshdadhich.com](mailto:me@harshdadhich.com) | (209) XXX-XXXX | [harshdadhich.com](https://harshdadhich.com) | [in hdadhich01](https://www.linkedin.com/in/hdadhich01) | [gh hdadhich01](https://github.com/hdadhich01)

## EDUCATION

### University of California, Los Angeles

Expected 2026

Master of Engineering: Data Science

Los Angeles, CA

- **Coursework:** Large-Scale Machine Learning, Big Data Analytics, Reinforcement Learning, Product Management

### University of California, Santa Cruz

Grad Spring 2025



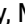


Bachelor of Science: Computer Science


Santa Cruz, CA

- **Involvement:** Google DSC, GitHub Ambassador, Delta Sigma Pi – **Awards:** Dean's Honors, Chancellor's Internship
- **Coursework:** Data Structures & Algorithms, Computer Systems & Architecture, Applied Machine Learning, Functional Prog, Computer Networks, Web Applications, Probability, Discrete Math, Linear Algebra

## SKILLS

**Languages:** Python, Java, C, C++, TypeScript, JavaScript, Bash, SQL, Haskell, HTML/CSS,  $\text{\LaTeX}$

**Frameworks & Tools:**  Django, FastAPI, Flask, Gunicorn, BeautifulSoup;  TensorFlow, PyTorch, scikit-learn, Pandas, NumPy, Matplotlib;  React.js, AngularJS, Next.js, Node.js, Bun, Tailwind;  Git, Jira, Linear;  Vercel, Postman, Figma

**Cloud & DB:** AWS (EC2, Lambda, Elastic Beanstalk, DynamoDB, S3), Google Cloud, Cloudflare (Workers, DNS, D1, KV), Docker, Linux (Debian), Terraform, Packer, Jenkins, jFrog, GitLab CI, GitHub Actions;  PostgreSQL, Chroma, Pinecone

**Concepts:** Agile, Networking (TCP/IP, DNS, HTTPS), REST APIs, Microservices, CI/CD, Exploratory Data Analysis, RAG

## WORK EXPERIENCE

### Dento AI | Software Engineer Intern

Jan 2025 – Jun 2025

FiDi-based startup providing an AI-powered platform for dental professionals

San Francisco, CA

- Optimized consultation transcription latency by **35%** with client-side audio streaming, achieving speeds  $\leq 5s$
- Achieved a **22x** server-side speedup with bulk transcription processing through **Deepgram** and multi-threading
- Implemented a schema + database partitioning script in **Django/PostgreSQL**, cutting query latency by **40%**

### Intuitive Surgical | DevOps Engineer Intern

Jun 2024 – Sep 2024

Industry leader in surgical robotics; AWS Operations Team under the Digital BU

Sunnyvale, CA

- Built reusable **Terraform** modules **reducing** setup time by **~90%** for AWS infra in multi-region environments
- Migrated CI/CD from Groovy/Java to GitLab CI YAML, streamlining security for an SDLC efficiency **increase of 60%**
- Trimmed manual release steps from **14**  $\rightarrow$  **2** with **Packer** + Terraform CI for AMI building and blue-green rollouts

### CITRIS and the Banatao Institute | Data Scientist Intern

May 2024 – Jul 2024

UC-wide research initiative; matched with host company: Beyond the Arc

Berkeley, CA


- Benchmarking and developing LLM tests on a nuclear energy project for an AI/ML Analytics B2B company
- Engineered text mining pipeline using OpenAI's **Assistants API**, with **~92%** accuracy for content classification
- Accelerated regulatory document processing by **40%** by parallelizing scraping and calls to the Assistants API

### UCSC Genomics Institute | Undergraduate Research Assistant

Apr 2024 – Dec 2024

Braingeneers Lab – experimentation with neural organoids; PI: Dr. Mircea Teodorescu

Santa Cruz, CA

- Mentored by Dr. Teodorescu to use **MEArec**, **matplotlib**, and **numpy** to simulate extracellular recordings
- Validated neural spike detection accuracy by comparing ground-truth neuron recordings with AI-sorted clusters, achieving **~95%** precision in identifying individual neuron activity patterns; more at  braingeneers/MEArec
- Formulated schemas for an autonomous agent framework to execute IoT-based wet lab experiments with LLMs

## PROJECTS

### **Movie Revenue Predictor – Neural Network Model Training** | Python, TensorFlow, Keras, Pandas

June 2024

- Built a movie revenue predictor boosting stock model accuracy by **~28%** using feed-forward networks
- Engineered features (director success, actor frequency) and robust data preprocessing pipeline for 5k+ movies

### **round-nutrition – FDA Compliant Nutritional Rounding Package** | PyPI, Docker, GitHub Actions

July 2022

- Shipped a package to round nutritional value data to meet federal standards and regulations (FDA, USDA, NDC)
- Received **20k+ downloads** from popular projects for use in nutritional interfaces and calculators

### **SlugTools – REST API for Aggregated University Data** | Flask, BeautifulSoup, Gunicorn, Cloudflare

May 2022

- Deployed a RESTful API on parsed/scraped university data for course catalog, dining, and facility services
- Streamlined deployment and improved server resource usage by **20%** with CI/CD and periodic scraping